practical-exam-06

May 23, 2023

[ ]:
```
from google.colab import drive
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).
```

1 Problem Statement 6

Create a Linear Regression Model using Python to predict home prices using the Boston Housing Dataset. The Boston Housing dataset contains information about various houses in Boston through different parameters. There are 506 samples and 14 feature variables in this dataset.

[ ]:
```
import pandas as pd
```

[160]:
```
data = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/exam_datasets/6.
↪housing_data.csv')
```

```
<IPython.core.display.HTML object>
```

[165]:
```
data.dropna()
```

```
<IPython.core.display.HTML object>
```

[165]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	
5	0.02985	0.0	2.18	0.0	0.458	6.430	58.7	6.0622	3	222	
..	…	…	…	…	…	…	…	…	…	…	
499	0.17783	0.0	9.69	0.0	0.585	5.569	73.5	2.3999	6	391	
500	0.22438	0.0	9.69	0.0	0.585	6.027	79.7	2.4982	6	391	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0

```
1       17.8  396.90   9.14  21.6
2       17.8  392.83   4.03  34.7
3       18.7  394.63   2.94  33.4
5       18.7  394.12   5.21  28.7
..       …     …    …   …
499     19.2  395.77  15.10  17.5
500     19.2  396.90  14.33  16.8
502     21.0  396.90   9.08  20.6
503     21.0  396.90   5.64  23.9
504     21.0  393.45   6.48  22.0

[394 rows x 14 columns]
```

[166]:
```
data.info()
```

```
<IPython.core.display.HTML object>
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column   Non-Null Count  Dtype
---  ------   --------------  -----
 0   CRIM     486 non-null    float64
 1   ZN       486 non-null    float64
 2   INDUS    486 non-null    float64
 3   CHAS     486 non-null    float64
 4   NOX      506 non-null    float64
 5   RM       506 non-null    float64
 6   AGE      486 non-null    float64
 7   DIS      506 non-null    float64
 8   RAD      506 non-null    int64
 9   TAX      506 non-null    int64
 10  PTRATIO  506 non-null    float64
 11  B        506 non-null    float64
 12  LSTAT    486 non-null    float64
 13  MEDV     506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

[173]:
```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score,␣
 ↪recall_score
```

```
<IPython.core.display.HTML object>
```

[180]:
```
data = data.dropna()
X = data.drop('MEDV', axis=1)
```

```
y = data['MEDV']


X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
lr = LinearRegression()
lr.fit(X_train, y_train)

y_pred = lr.predict(X_test)
```

<IPython.core.display.HTML object>

[181]:
```
r2 = lr.score(X_test, y_test)
print(f'R^2 score: {r2:.2f}')
```

<IPython.core.display.HTML object>

R^2 score: 0.79

[182]:
```
import matplotlib.pyplot as plt

# Plot the predicted and actual home prices
plt.scatter(y_test, y_pred)
plt.xlabel('Actual Home Prices')
plt.ylabel('Predicted Home Prices')
plt.title('Actual vs Predicted Home Prices')
plt.show()
```

<IPython.core.display.HTML object>

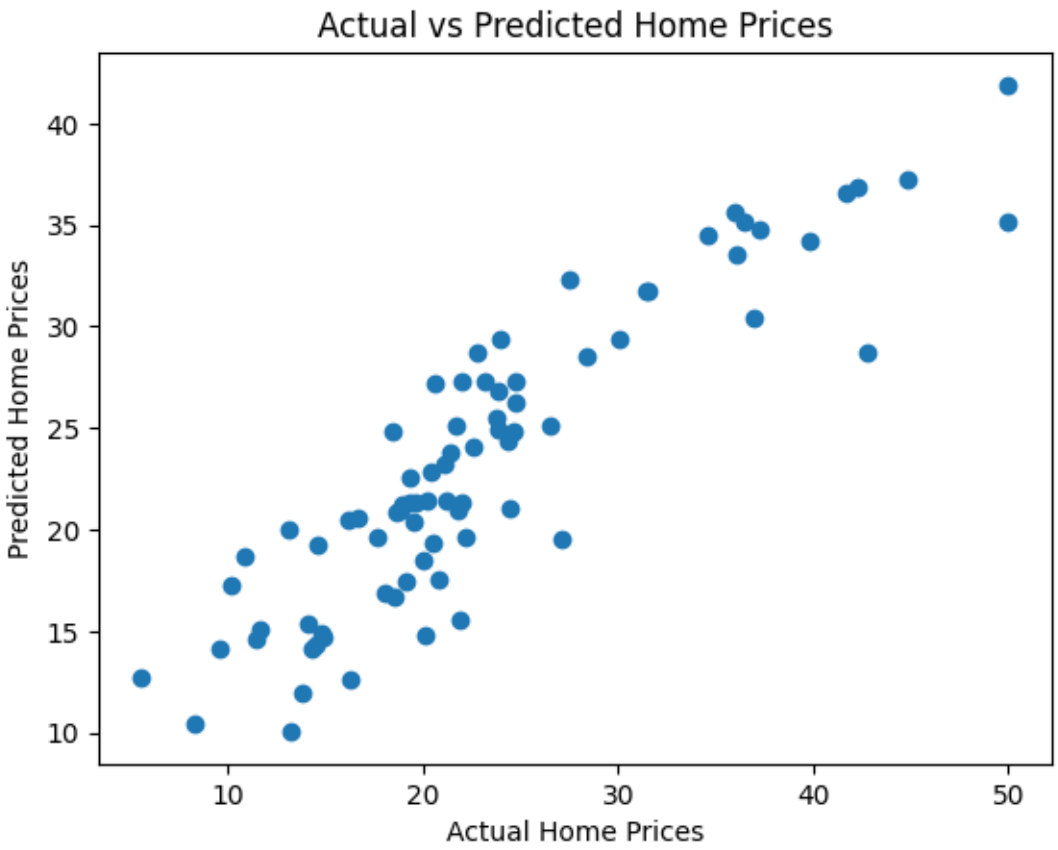
Actual vs Predicted Home Prices
Predicted Home Prices
Actual Home Prices
10
15
20
25
30
35
40
10
20
30
40
50